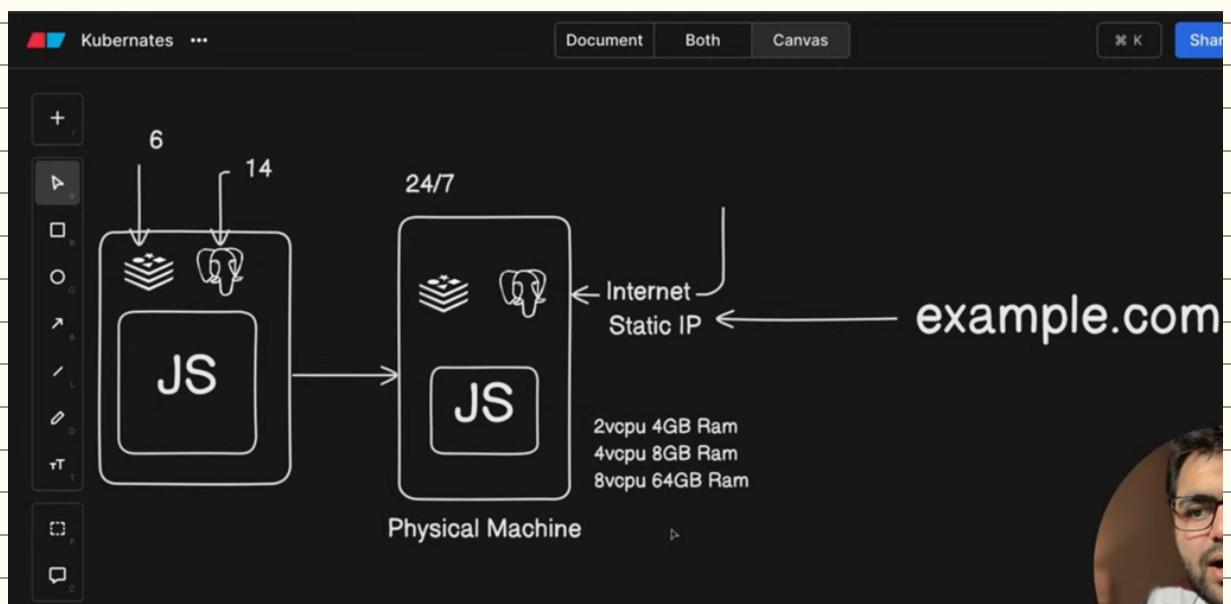
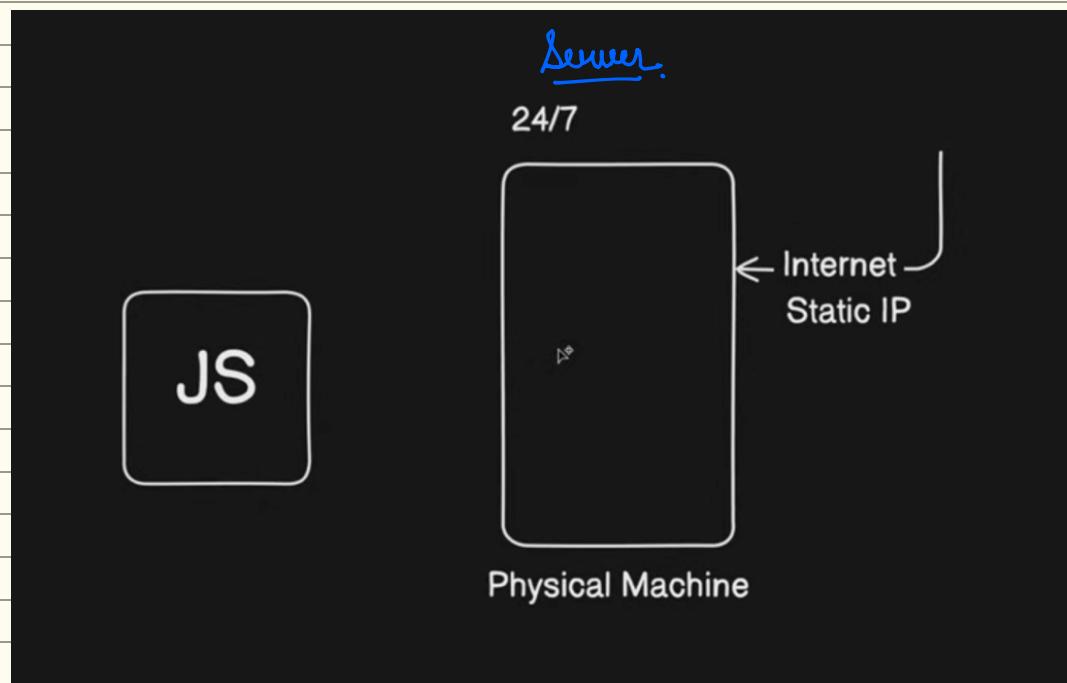
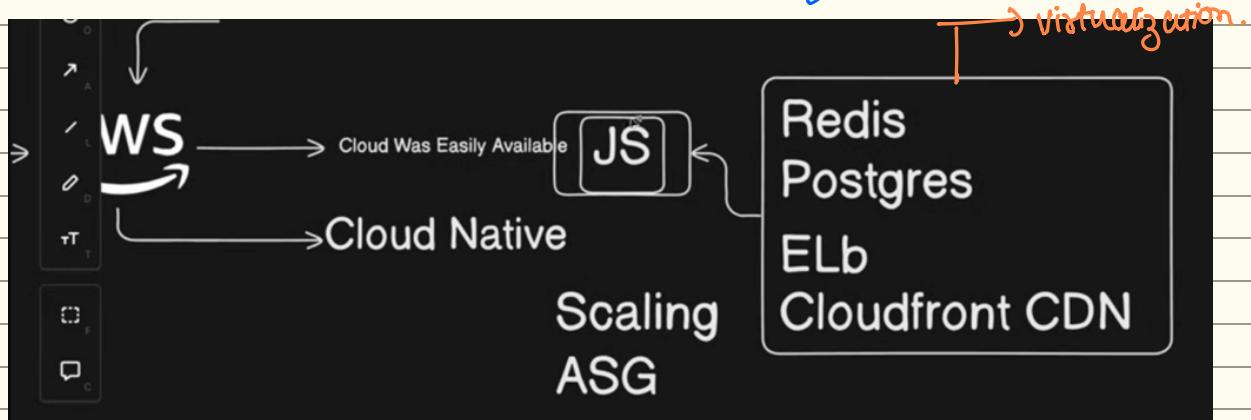


KUBERNATES

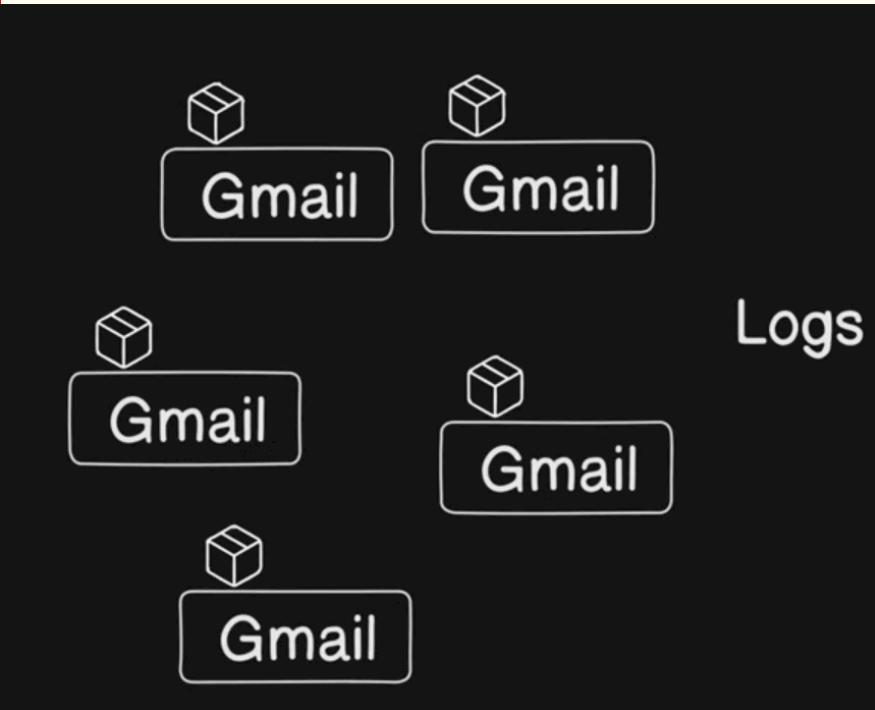
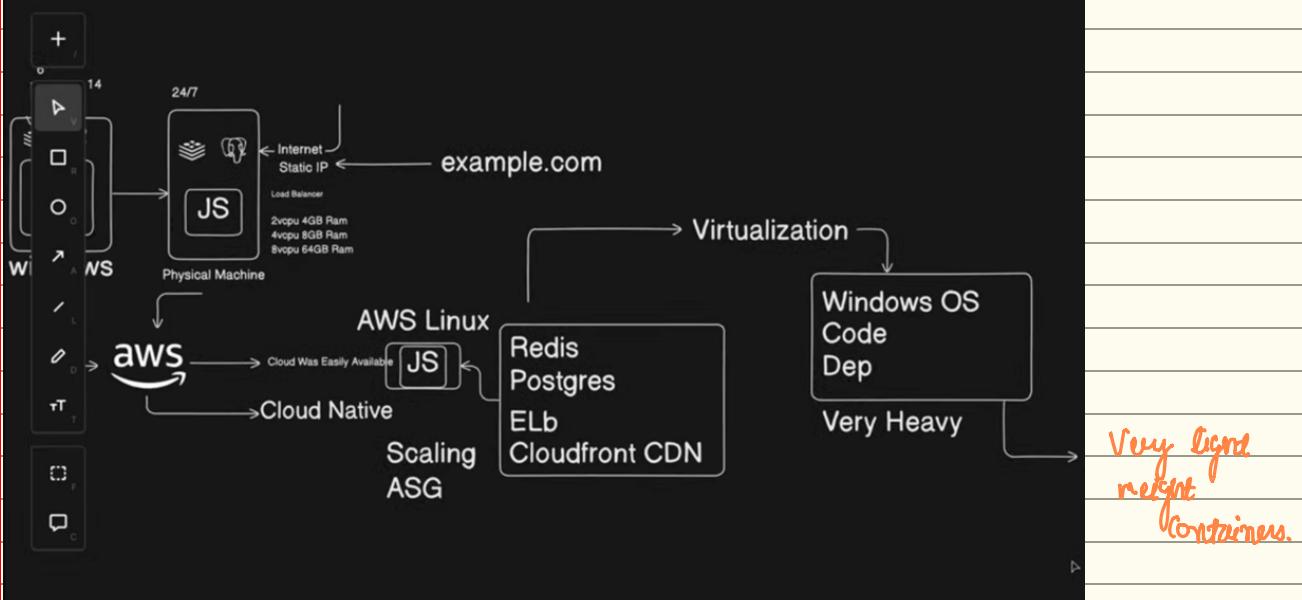
↳ Explained !!!



After the emergence of AWS:-



Common Issue:- It works on my computer, hence people started migrating to cloud Native tech, where virtualization is reqd.



→ Container Orchestration:-

Container orchestration is the process of automating the management of containerized applications, including their deployment, scaling, and security. It helps organizations to take advantage of containerization without incurring extra maintenance costs.

∴ In the early era of Containerization,

↳ Google was one of the early staged company, adopting to container orchestration.

∴ It introduced Borg, for internal container management.

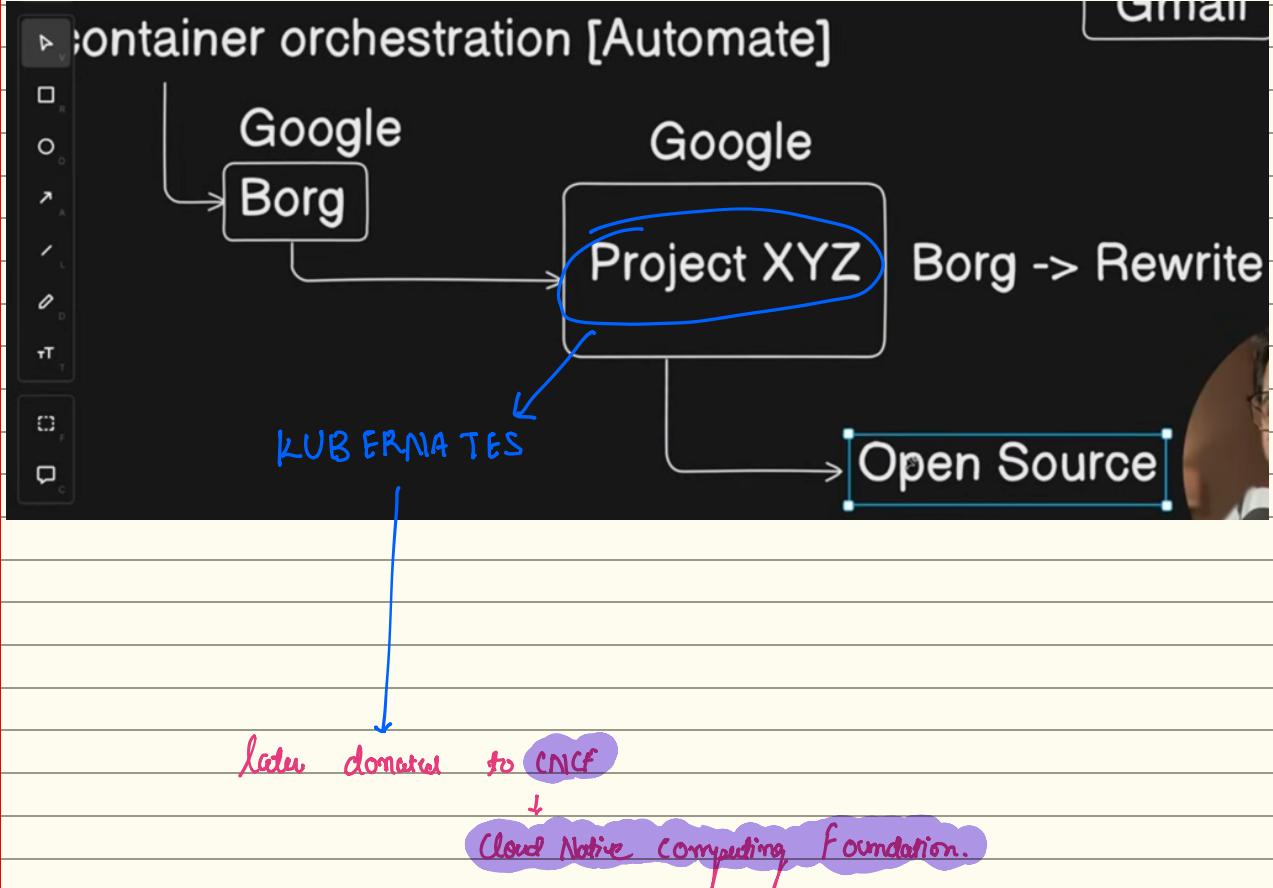


Google Research

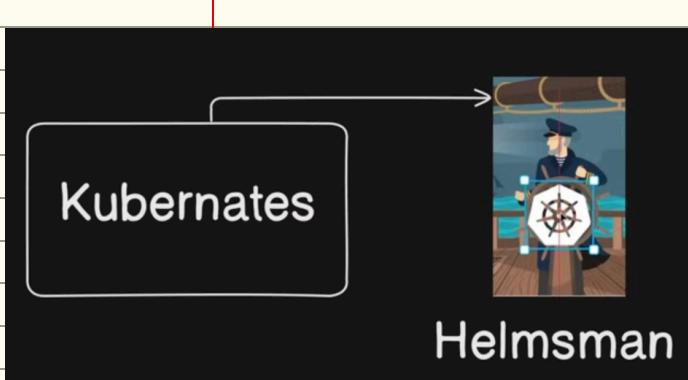
<http://research.google/pubs/large-scale-cluster-manag...> ::

Large-scale cluster management at {Google} with {Borg}

Google's Borg system is a cluster manager that runs hundreds of thousands of jobs, from many thousands of different applications, across a number of ...



→ Kubernetes:-



Kubernetes (often abbreviated as K8s) is an open-source container orchestration platform used to automate:

- ✓ Deployment of containerized applications
- ✓ Scaling (adding/removing instances automatically)
- ✓ Load balancing
- ✓ Resource management
- ✓ Rolling updates and rollbacks
- ✓ Self-healing (restart/replace failed containers)

Why do we need Kubernetes?

When you have **many applications** running in containers (like Docker), managing them manually gets complex:

Without Kubernetes

Manual deployment of containers

Hard to recover failures

Hard to distribute load

With Kubernetes

Automated deployment and scaling

Self-healing automatically

Built-in load balancing

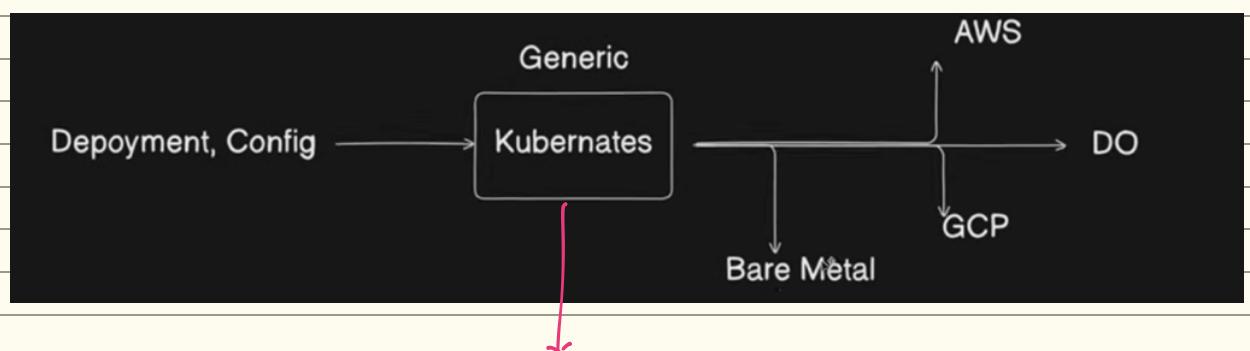
Kubernetes, also known as K8s, is an open source system for automating deployment, scaling, and management of containerized applications.

It groups containers that make up an application into logical units for easy management and discovery. Kubernetes builds upon **15 years of experience of running production workloads at Google**, combined with best-of-breed ideas and practices from the community.

| Key Concepts | |
|--------------|--|
| Term | Meaning |
| Cluster | A group of machines (nodes) managed by K8s |
| Node | A worker machine (VM or physical server) |
| Pod | Smallest deployable unit → one or more containers |
| Deployment | Blueprint for how many pods, what version, scaling rules |
| Service | Stable network endpoint to access pods |
| Ingress | Routes external traffic to internal services |

Example Use Cases

- Microservices applications
- Scalable web platforms (Netflix, Spotify use it)
- CI/CD pipelines
- Edge / hybrid cloud infrastructure



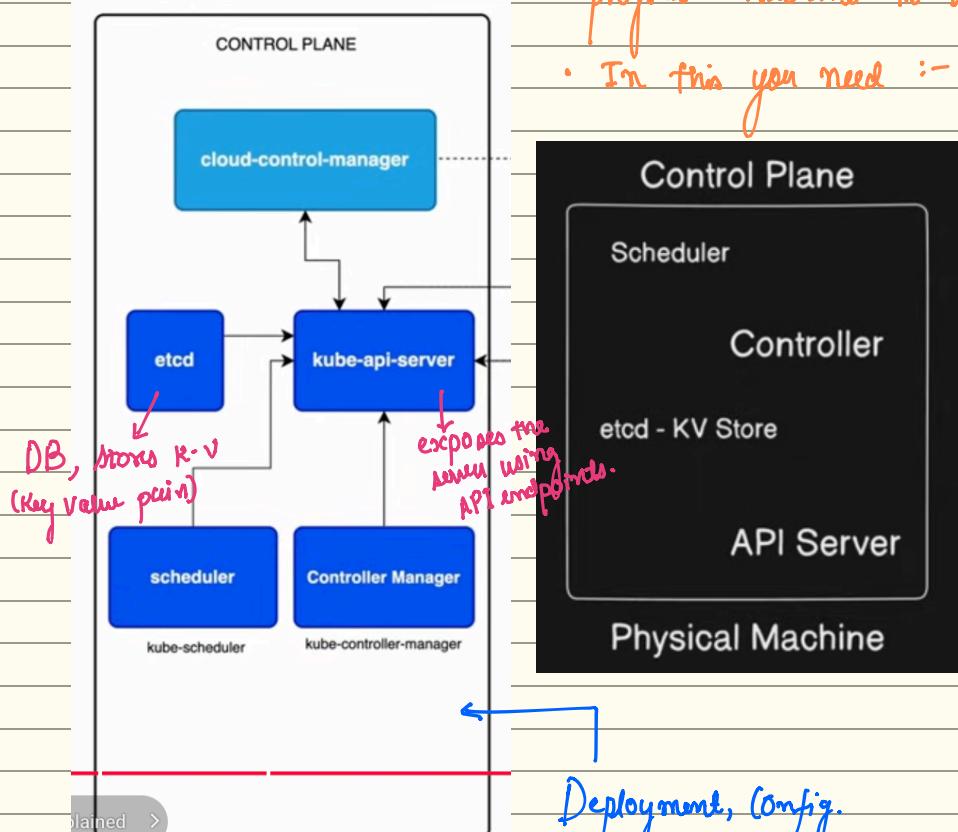
Helps to develop generic configuration for deployment on all cloud native platforms, hence no vendor locking.

→ **Architecture for Kubernetes :-**

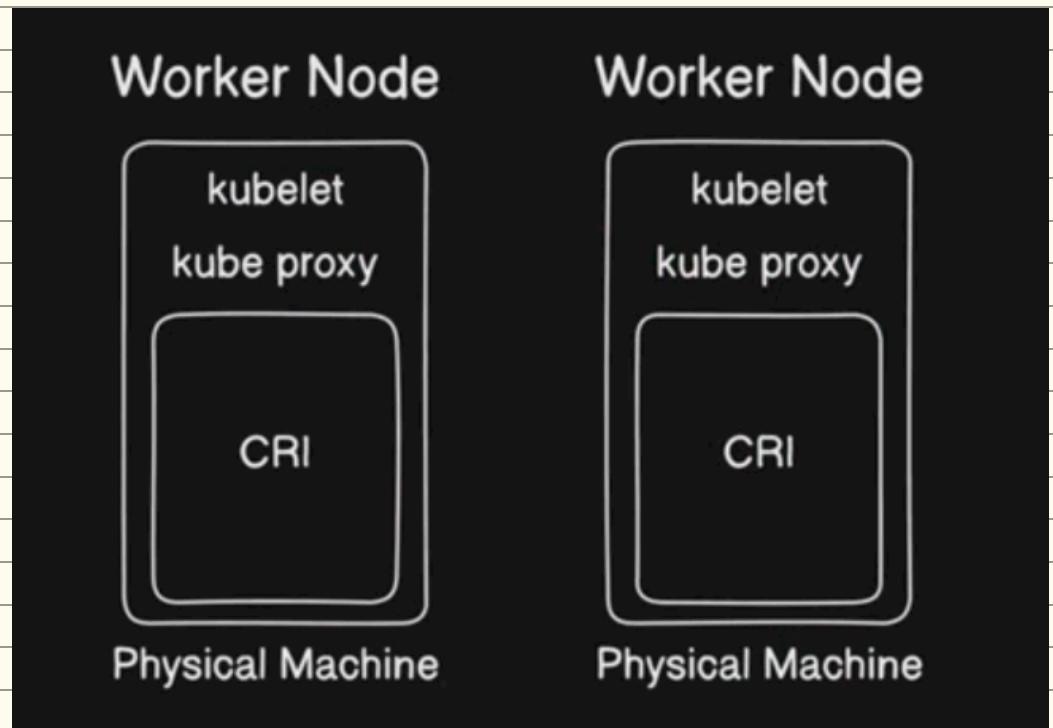
• physical machine is the control plane.

• In this you need :-

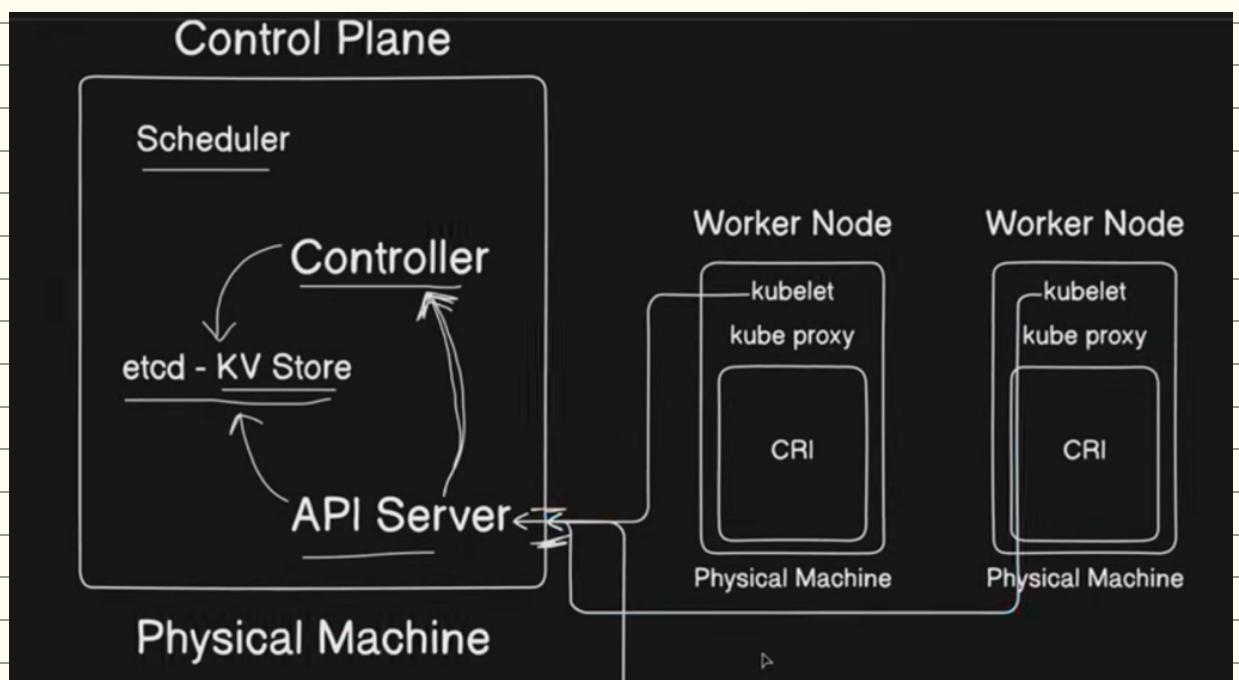
- (i). Scheduler
- (ii). Controller
- (iii). etcd - KV Store
- (iv). API Server.

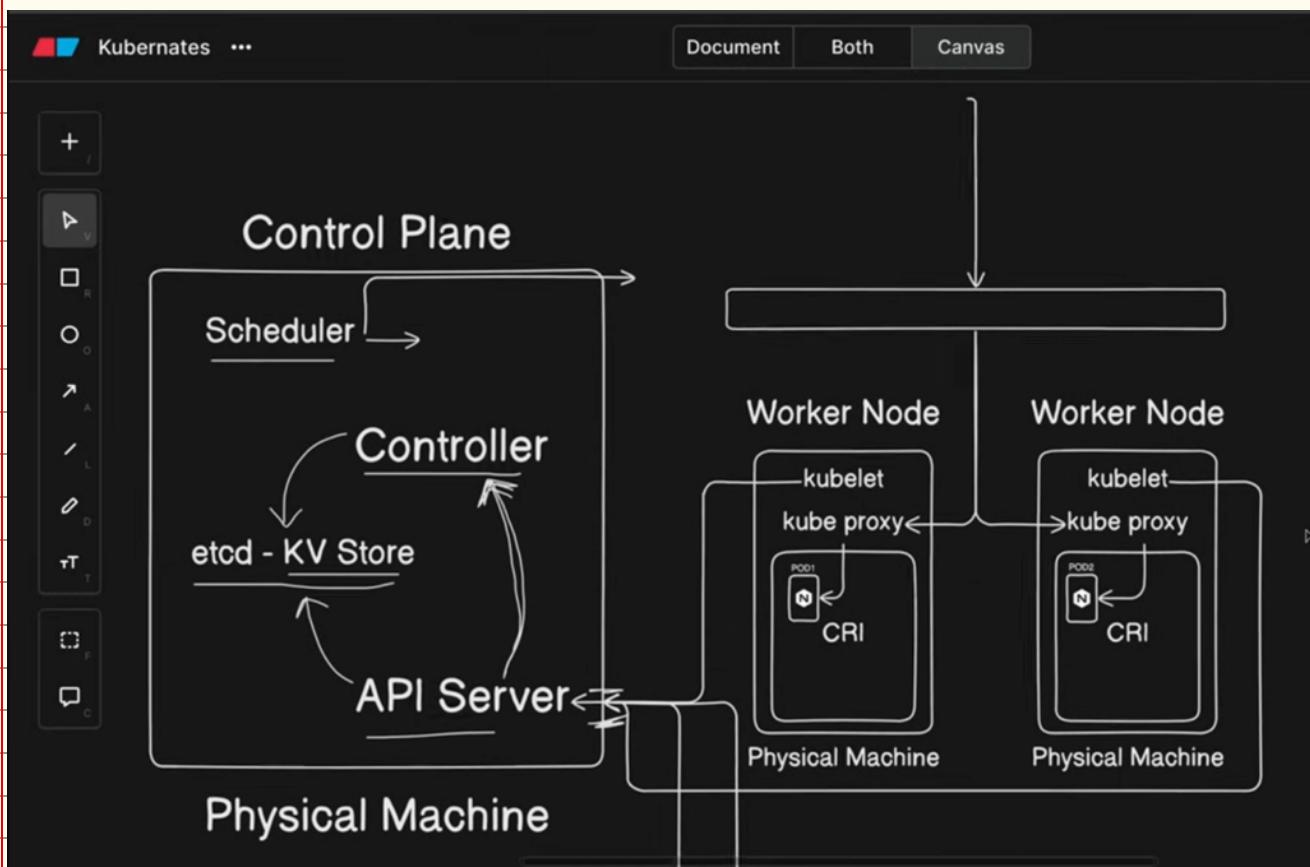
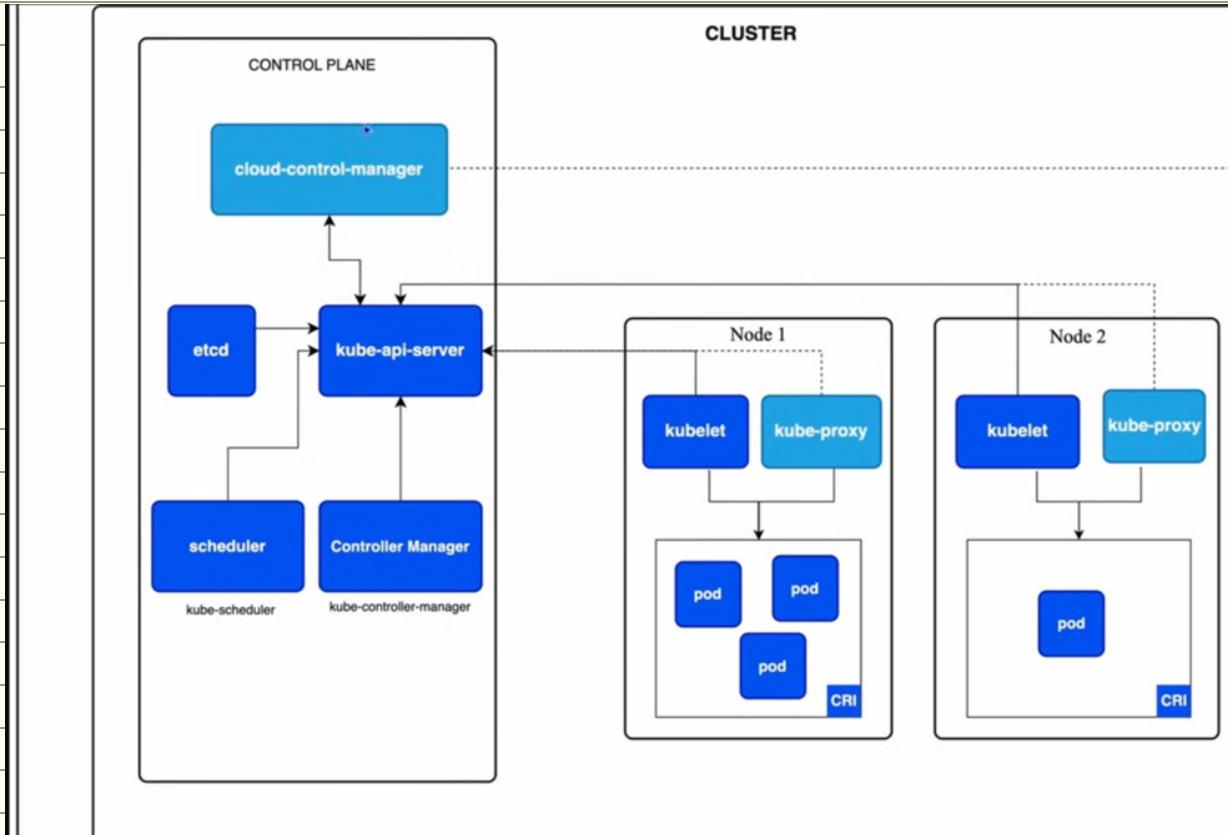


- Worker Node are different physical machines, where the work load runs.
 - ↳ Here the actual (containers) applications run on this worker nodes.



• The kubelets in the worker nodes are connected to the API server.





• **kube proxy** provides:- Networking and Traffic rules to the **kubelets**.

→ Container Runtimes :-

Container Runtimes

1. Podman
2. CRI-O
3. Containerd (CNCF)
4. LXC (Linux Containers)
5. Kata Containers
6. Firecracker (AWS Lambda)

→ Working of Kubernetes :-

🚀 Kubernetes Cluster — High-Level Working

A Kubernetes cluster consists of:

1 Control Plane (Brain of the system)

Manages what should run and where.

2 Worker Nodes (Muscles of the system)

These actually run the application containers.

🧠 Control Plane Components

| Component | Responsibility |
|--------------------|---|
| API Server | Entry point — all communication happens through this (kubectl → API Server → Cluster) |
| Scheduler | Decides which node will run each Pod based on resources |
| Controller Manager | Ensures the desired state: if a pod crashes → recreate |
| etcd | Key-value store — keeps the cluster state (like a database for configs) |

💡 Worker Node Components

| Component | Responsibility |
|-------------------|---|
| Kubelet | Talks to the API Server, starts/stops containers on that node |
| Container Runtime | Runs containers (Docker, containerd, CRI-O) |
| Kube-proxy | Handles networking & load balancing to Pods |

📌 How Kubernetes Works (Step-by-Step)

Example: You tell Kubernetes to run 5 instances of your app.

1. 🧑 You submit a Deployment YAML → `kubectl apply -f app.yaml`
2. 🖥 API Server receives request → stores desired state in `etcd`
3. 🧠 Scheduler checks all nodes → picks the best nodes for 5 Pods
4. 💡 Kubelet on selected nodes runs containers via container runtime
5. 📱 Controller Manager keeps checking:
 - If any Pod crashes → recreate automatically
 - If demand increases → scale up
 - If node dies → move Pods to another node
6. 🌐 kube-proxy ensures networking between Pods and Services

Cluster Working Visualization

lua

Copy code

